

dlnd_face_generation

July 11, 2021

1 Face Generation

In this project, you'll define and train a DCGAN on a dataset of faces. Your goal is to get a generator network to generate *new* images of faces that look as realistic as possible!

The project will be broken down into a series of tasks from **loading in data to defining and training adversarial networks**. At the end of the notebook, you'll be able to visualize the results of your trained Generator to see how it performs; your generated samples should look like fairly realistic faces with small amounts of noise.

1.0.1 Get the Data

You'll be using the [CelebFaces Attributes Dataset \(CelebA\)](#) to train your adversarial networks.

This dataset is more complex than the number datasets (like MNIST or SVHN) you've been working with, and so, you should prepare to define deeper networks and train them for a longer time to get good results. It is suggested that you utilize a GPU for training.

1.0.2 Pre-processed Data

Since the project's main focus is on building the GANs, we've done *some* of the pre-processing for you. Each of the CelebA images has been cropped to remove parts of the image that don't include a face, then resized down to 64x64x3 NumPy images. Some sample data is show below.

If you are working locally, you can download this data [by clicking here](#)

This is a zip file that you'll need to extract in the home directory of this notebook for further loading and processing. After extracting the data, you should be left with a directory of data `processed_celeba_small/`

```
In [1]: # can comment out after executing
        !unzip processed_celeba_small.zip
```

```
Archive:  processed_celeba_small.zip
replace processed_celeba_small/.DS_Store? [y]es, [n]o, [A]ll, [N]one, [r]ename: ^C
```

```
In [1]: data_dir = 'processed_celeba_small/'
```

```
"""
```

```

DON'T MODIFY ANYTHING IN THIS CELL
"""
import pickle as pkl
import matplotlib.pyplot as plt
import numpy as np
import problem_unittests as tests
#import helper

%matplotlib inline

```

1.1 Visualize the CelebA Data

The [CelebA](#) dataset contains over 200,000 celebrity images with annotations. Since you're going to be generating faces, you won't need the annotations, you'll only need the images. Note that these are color images with 3 color channels (RGB) each.

1.1.1 Pre-process and Load the Data

Since the project's main focus is on building the GANs, we've done *some* of the pre-processing for you. Each of the CelebA images has been cropped to remove parts of the image that don't include a face, then resized down to 64x64x3 NumPy images. This *pre-processed* dataset is a smaller subset of the very large CelebA data.

There are a few other steps that you'll need to **transform** this data and create a **DataLoader**.

Exercise: Complete the following `get_dataloader` function, such that it satisfies these requirements:

- Your images should be square, Tensor images of size `image_size x image_size` in the x and y dimension.
- Your function should return a `DataLoader` that shuffles and batches these Tensor images.

ImageFolder To create a dataset given a directory of images, it's recommended that you use PyTorch's [ImageFolder](#) wrapper, with a root directory `processed_celeba_small/` and data transformation passed in.

```

In [2]: # necessary imports
import torch
from torchvision import datasets
from torchvision import transforms

In [3]: def get_dataloader(batch_size, image_size, data_dir='processed_celeba_small/'):
        """
        Batch the neural network data using DataLoader
        :param batch_size: The size of each batch; the number of images in a batch
        :param img_size: The square size of the image data (x, y)
        :param data_dir: Directory where image data is located
        :return: DataLoader with batched data

```

```

"""

# TODO: Implement function and return a dataloader
transform = transforms.Compose([transforms.Resize(image_size),
                                transforms.ToTensor()])
train_data = datasets.ImageFolder(data_dir, transform=transform)

return torch.utils.data.DataLoader(train_data, batch_size=batch_size)

```

1.2 Create a DataLoader

Exercise: Create a DataLoader `celeba_train_loader` **with appropriate hyperparameters**. Call the above function and create a dataloader to view images. * You can decide on any reasonable `batch_size` parameter * Your `image_size` **must be 32**. Resizing the data to a smaller size will make for faster training, while still creating convincing images of faces!

```

In [4]: # Define function hyperparameters
        batch_size = 128
        img_size = 32

        """
        DON'T MODIFY ANYTHING IN THIS CELL THAT IS BELOW THIS LINE
        """

        # Call your function and get a dataloader
        celeba_train_loader = get_dataloader(batch_size, img_size)

```

Next, you can view some images! You should see square images of somewhat-centered faces.

Note: You'll need to convert the Tensor images into a NumPy type and transpose the dimensions to correctly display an image, suggested `imshow` code is below, but it may not be perfect.

```


In [5]: # helper display function
        def imshow(img):
            npimg = img.numpy()
            plt.imshow(np.transpose(npimg, (1, 2, 0)))

        """
        DON'T MODIFY ANYTHING IN THIS CELL THAT IS BELOW THIS LINE
        """

        # obtain one batch of training images
        dataiter = iter(celeba_train_loader)
        images, _ = dataiter.next() # _ for no labels

        # plot the images in the batch, along with the corresponding labels
        fig = plt.figure(figsize=(20, 4))
        plot_size=20
        for idx in np.arange(plot_size):
            ax = fig.add_subplot(2, plot_size/2, idx+1, xticks=[], yticks=[])
            imshow(images[idx])

```



Exercise: Pre-process your image data and scale it to a pixel range of -1 to 1 You need to do a bit of pre-processing; you know that the output of a tanh activated generator will contain pixel values in a range from -1 to 1, and so, we need to rescale our training images to a range of -1 to 1. (Right now, they are in a range from 0-1.)

```
In [6]: # TODO: Complete the scale function
def scale(x, feature_range=(-1, 1)):
    ''' Scale takes in an image x and returns that image, scaled
        with a feature_range of pixel values from -1 to 1.
        This function assumes that the input x is already scaled from 0-1. '''
    # assume x is scaled to (0, 1)
    # scale to feature_range and return scaled x
    low, high = feature_range
    x = x*(high-low)+low
    return x
```

```
In [7]: """
DON'T MODIFY ANYTHING IN THIS CELL THAT IS BELOW THIS LINE
"""

# check scaled range
# should be close to -1 to 1
img = images[0]
scaled_img = scale(img)

print('Min: ', scaled_img.min())
print('Max: ', scaled_img.max())
```

```
Min: tensor(-0.9529)
Max: tensor(0.9451)
```

2 Define the Model

A GAN is comprised of two adversarial networks, a discriminator and a generator.

2.1 Discriminator

Your first task will be to define the discriminator. This is a convolutional classifier like you've built before, only without any maxpooling layers. To deal with this complex data, it's suggested you use a deep network with **normalization**. You are also allowed to create any helper functions that may be useful.

Exercise: Complete the Discriminator class

- The inputs to the discriminator are 32x32x3 tensor images
- The output should be a single value that will indicate whether a given image is real or fake

```
In [8]: import torch.nn as nn
import torch.nn.functional as F
```

```
In [9]: class Discriminator(nn.Module):
    #ill be making this architecture:
    #https://raw.githubusercontent.com/udacity/deep-learning-v2-pytorch/b82f18222e46c271
    def __init__(self, conv_dim):
        """
        Initialize the Discriminator Module
        :param conv_dim: The depth of the first convolutional layer
        """
        super(Discriminator, self).__init__()
        self.conv_dim = conv_dim
        # complete init function
        self.conv1 = nn.Conv2d(3, conv_dim, 4, stride=2, padding=1)
        # sees 16x16x..
        self.conv2 = nn.Conv2d(conv_dim, conv_dim*2, 4, stride=2, padding=1)
        # sees 8x8x..
        self.batch1 = nn.BatchNorm2d(conv_dim*2)
        self.conv3 = nn.Conv2d(conv_dim*2, conv_dim*4, 4, stride=2, padding=1)
        self.batch2 = nn.BatchNorm2d(conv_dim*4)
        # sees 4x4x..
        self.fc = nn.Linear(4*4*conv_dim*4, 1)

    def forward(self, x):
        """
        Forward propagation of the neural network
        :param x: The input to the neural network
        :return: Discriminator logits; the output of the neural network
        """
        # define feedforward behavior
        x = F.leaky_relu(self.conv1(x), 0.2)
        x = self.batch1(F.leaky_relu(self.conv2(x), 0.2))
        x = self.batch2(F.leaky_relu(self.conv3(x), 0.2))
        x = x.view(-1, 4*4*self.conv_dim*4)
        x = self.fc(x)
```

```

        return x

    """
    DON'T MODIFY ANYTHING IN THIS CELL THAT IS BELOW THIS LINE
    """
    tests.test_discriminator(Discriminator)

```

Tests Passed

2.2 Generator

The generator should upsample an input and generate a *new* image of the same size as our training data 32x32x3. This should be mostly transpose convolutional layers with normalization applied to the outputs.

Exercise: Complete the Generator class

- The inputs to the generator are vectors of some length `z_size`
- The output should be a image of shape 32x32x3

In [10]: `class Generator(nn.Module):`

```

    def __init__(self, z_size, conv_dim):
        #ill be making this network
        #https://github.com/udacity/deep-learning-v2-pytorch/raw/b82f18222e46c27138fa14
        """
        Initialize the Generator Module
        :param z_size: The length of the input latent vector, z
        :param conv_dim: The depth of the inputs to the *last* transpose convolutional
        """

        super(Generator, self).__init__()

        # complete init function
        self.z_size = z_size
        self.conv_dim = conv_dim
        self.fc = nn.Linear(z_size, 4*4*conv_dim*4)
        self.conv1 = nn.ConvTranspose2d(conv_dim*4, conv_dim*2, 4, stride=2, padding=1, b
        self.batch1 = nn.BatchNorm2d(conv_dim*2)
        self.conv2 = nn.ConvTranspose2d(conv_dim*2, conv_dim, 4, stride=2, padding=1, b
        self.batch2 = nn.BatchNorm2d(conv_dim)
        self.conv3 = nn.ConvTranspose2d(conv_dim, 3, 4, stride=2, padding=1, bias=False

    def forward(self, x):
        """
        Forward propagation of the neural network
        :param x: The input to the neural network

```

```

        :return: A 32x32x3 Tensor image as output
        """
        # define feedforward behavior
        x = self.fc(x)
        x = x.view(-1, self.conv_dim*4, 4, 4) # (batch_size, depth, 4, 4)
        x = self.batch1(F.relu(self.conv1(x)))
        x = self.batch2(F.relu(self.conv2(x)))
        x = F.tanh(self.conv3(x))
        return x

    """
    DON'T MODIFY ANYTHING IN THIS CELL THAT IS BELOW THIS LINE
    """
    tests.test_generator(Generator)

```

Tests Passed

2.3 Initialize the weights of your networks

To help your models converge, you should initialize the weights of the convolutional and linear layers in your model. From reading the [original DCGAN paper](#), they say: > All weights were initialized from a zero-centered Normal distribution with standard deviation 0.02.

So, your next task will be to define a weight initialization function that does just this!

You can refer back to the lesson on weight initialization or even consult existing model code, such as that from [the networks.py file in CycleGAN Github repository](#) to help you complete this function.

Exercise: Complete the weight initialization function

- This should initialize only **convolutional** and **linear** layers
- Initialize the weights to a normal distribution, centered around 0, with a standard deviation of 0.02.
- The bias terms, if they exist, may be left alone or set to 0.

```

In [11]: from torch.nn import init
         def weights_init_normal(m):
             """
             Applies initial weights to certain layers in a model .
             The weights are taken from a normal distribution
             with mean = 0, std dev = 0.02.
             :param m: A module or layer in a network
             """
             # classname will be something like:
             # `Conv`, `BatchNorm2d`, `Linear`, etc.
             classname = m.__class__.__name__
             # TODO: Apply initial weights to convolutional and linear layers
             if hasattr(m, 'weight') and (classname.find('Conv') != -1 or classname.find('Linear')

```

```

init.normal_(m.weight.data, 0.0, 0.02)
if hasattr(m, 'bias') and m.bias is not None:
    init.constant_(m.bias.data, 0.0)

```

2.4 Build complete network

Define your models' hyperparameters and instantiate the discriminator and generator from the classes defined above. Make sure you've passed in the correct input arguments.

```

In [12]: """
DON'T MODIFY ANYTHING IN THIS CELL THAT IS BELOW THIS LINE
"""

def build_network(d_conv_dim, g_conv_dim, z_size):
    # define discriminator and generator
    D = Discriminator(d_conv_dim)
    G = Generator(z_size=z_size, conv_dim=g_conv_dim)

    # initialize model weights
    D.apply(weights_init_normal)
    G.apply(weights_init_normal)

    print(D)
    print()
    print(G)

    return D, G

```

Exercise: Define model hyperparameters

```

In [13]: # Define model hyperparams
d_conv_dim = 32
g_conv_dim = 32
z_size = 100

"""
DON'T MODIFY ANYTHING IN THIS CELL THAT IS BELOW THIS LINE
"""

D, G = build_network(d_conv_dim, g_conv_dim, z_size)

```

```

Discriminator(
  (conv1): Conv2d(3, 32, kernel_size=(4, 4), stride=(2, 2), padding=(1, 1))
  (conv2): Conv2d(32, 64, kernel_size=(4, 4), stride=(2, 2), padding=(1, 1))
  (batch1): BatchNorm2d(64, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
  (conv3): Conv2d(64, 128, kernel_size=(4, 4), stride=(2, 2), padding=(1, 1))
  (batch2): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
  (fc): Linear(in_features=2048, out_features=1, bias=True)
)

```



```

Generator(
  (fc): Linear(in_features=100, out_features=2048, bias=True)
  (conv1): ConvTranspose2d(128, 64, kernel_size=(4, 4), stride=(2, 2), padding=(1, 1), bias=False)
  (batch1): BatchNorm2d(64, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
  (conv2): ConvTranspose2d(64, 32, kernel_size=(4, 4), stride=(2, 2), padding=(1, 1), bias=False)
  (batch2): BatchNorm2d(32, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
  (conv3): ConvTranspose2d(32, 3, kernel_size=(4, 4), stride=(2, 2), padding=(1, 1), bias=False)
)

```

2.4.1 Training on GPU

Check if you can train on GPU. Here, we'll set this as a boolean variable `train_on_gpu`. Later, you'll be responsible for making sure that `> * Models, * Model inputs, and * Loss function arguments`

Are moved to GPU, where appropriate.

```

In [14]: """
         DON'T MODIFY ANYTHING IN THIS CELL
         """
         import torch

         # Check for a GPU
         train_on_gpu = torch.cuda.is_available()
         if not train_on_gpu:
             print('No GPU found. Please use a GPU to train your neural network.')
         else:
             print('Training on GPU!')

```

Training on GPU!

2.5 Discriminator and Generator Losses

Now we need to calculate the losses for both types of adversarial networks.

2.5.1 Discriminator Losses

- For the discriminator, the total loss is the sum of the losses for real and fake images, $d_loss = d_real_loss + d_fake_loss$.
- Remember that we want the discriminator to output 1 for real images and 0 for fake images, so we need to set up the losses to reflect that.

2.5.2 Generator Loss

The generator loss will look similar only with flipped labels. The generator's goal is to get the discriminator to *think* its generated images are *real*.

Exercise: Complete real and fake loss functions You may choose to use either cross entropy or a least squares error loss to complete the following `real_loss` and `fake_loss` functions.

```
In [15]: def real_loss(D_out):
    '''Calculates how close discriminator outputs are to being real.
        param, D_out: discriminator logits
        return: real loss'''
    labels = torch.ones_like(D_out.squeeze())
    if train_on_gpu:
        labels = labels.cuda()
    criterion = nn.BCEWithLogitsLoss()
    loss = criterion(D_out.squeeze(), labels)
    return loss

def fake_loss(D_out):
    '''Calculates how close discriminator outputs are to being fake.
        param, D_out: discriminator logits
        return: fake loss'''
    labels = torch.zeros_like(D_out.squeeze())
    if train_on_gpu:
        labels = labels.cuda()
    criterion = nn.BCEWithLogitsLoss()
    loss = criterion(D_out.squeeze(), labels)
    return loss
```

2.6 Optimizers

Exercise: Define optimizers for your Discriminator (D) and Generator (G) Define optimizers for your models with appropriate hyperparameters.

```
In [16]: import torch.optim as optim

#following https://arxiv.org/pdf/1511.06434.pdf

# params
learning_rate = 0.0002
beta_1 = 0.5
beta_2 = 0.999

# Create optimizers for the discriminator D and generator G
d_optimizer = optim.Adam(D.parameters(), learning_rate, [beta_1, beta_2])
g_optimizer = optim.Adam(G.parameters(), learning_rate, [beta_1, beta_2])
```

2.7 Training

Training will involve alternating between training the discriminator and the generator. You'll use your functions `real_loss` and `fake_loss` to help you calculate the discriminator losses.

- You should train the discriminator by alternating on real and fake images
- Then the generator, which tries to trick the discriminator and should have an opposing loss function

Saving Samples You've been given some code to print out some loss statistics and save some generated "fake" samples.

Exercise: Complete the training function Keep in mind that, if you've moved your models to GPU, you'll also have to move any model inputs to GPU.

```
In [17]: def train(D, G, n_epochs, print_every=50):
    '''Trains adversarial networks for some number of epochs
    param, D: the discriminator network
    param, G: the generator network
    param, n_epochs: number of epochs to train for
    param, print_every: when to print and record the models' losses
    return: D and G losses'''

    # move models to GPU
    if train_on_gpu:
        D.cuda()
        G.cuda()

    # keep track of loss and generated, "fake" samples
    samples = []
    losses = []

    # Get some fixed data for sampling. These are images that are held
    # constant throughout training, and allow us to inspect the model's performance
    sample_size=16
    fixed_z = np.random.uniform(-1, 1, size=(sample_size, z_size))
    fixed_z = torch.from_numpy(fixed_z).float()
    # move z to GPU if available
    if train_on_gpu:
        fixed_z = fixed_z.cuda()

    # epoch training loop
    for epoch in range(n_epochs):

        # batch training loop
        for batch_i, (real_images, _) in enumerate(celeba_train_loader):

            batch_size = real_images.size(0)
            real_images = scale(real_images)

            # =====
            #          YOUR CODE HERE: TRAIN THE NETWORKS
```

```

# =====

# 1. Train the discriminator on real and fake images
d_optimizer.zero_grad()
if train_on_gpu:
    real_images = real_images.cuda()
D_real = D(real_images)
d_real_loss = real_loss(D_real)
z = np.random.uniform(-1, 1, size=(batch_size, z_size))
z = torch.from_numpy(z).float()
if train_on_gpu:
    z = z.cuda()
fake_images = G(z)
D_fake = D(fake_images)
d_fake_loss = fake_loss(D_fake)
d_loss = d_real_loss + d_fake_loss
d_loss.backward()
d_optimizer.step()

# 2. Train the generator with an adversarial loss
g_optimizer.zero_grad()
z = np.random.uniform(-1, 1, size=(batch_size, z_size))
z = torch.from_numpy(z).float()
if train_on_gpu:
    z = z.cuda()
fake_images = G(z)
D_fake = D(fake_images)
g_loss = real_loss(D_fake)
g_loss.backward()
g_optimizer.step()

# =====
#                               END OF YOUR CODE
# =====

# Print some loss stats
if batch_i % print_every == 0:
    # append discriminator loss and generator loss
    losses.append((d_loss.item(), g_loss.item()))
    # print discriminator and generator loss
    print('Epoch [{:5d}/{:5d}] | d_loss: {:.4f} | g_loss: {:.4f}'.format(
        epoch+1, n_epochs, d_loss.item(), g_loss.item()))

## AFTER EACH EPOCH##
# this code assumes your generator is named G, feel free to change the name
# generate and save sample, fake images

```

```

        G.eval() # for generating samples
        samples_z = G(fixed_z)
        samples.append(samples_z)
        G.train() # back to training mode

    # Save training generator samples
    with open('train_samples.pkl', 'wb') as f:
        pickle.dump(samples, f)

    # finally return losses
    return losses

```

Set your number of training epochs and train your GAN!

```

In [18]: # set number of epochs
         n_epochs = 20

         """
         DON'T MODIFY ANYTHING IN THIS CELL
         """

         # call training function
         losses = train(D, G, n_epochs=n_epochs)

```

Epoch [1/	20]		d_loss: 1.7427		g_loss: 0.5970
Epoch [1/	20]		d_loss: 0.0740		g_loss: 3.6904
Epoch [1/	20]		d_loss: 0.0449		g_loss: 4.2517
Epoch [1/	20]		d_loss: 0.1251		g_loss: 3.5460
Epoch [1/	20]		d_loss: 0.5474		g_loss: 2.5696
Epoch [1/	20]		d_loss: 0.7517		g_loss: 1.5036
Epoch [1/	20]		d_loss: 0.8189		g_loss: 1.2194
Epoch [1/	20]		d_loss: 0.7837		g_loss: 1.9027
Epoch [1/	20]		d_loss: 0.9164		g_loss: 1.1374
Epoch [1/	20]		d_loss: 0.8254		g_loss: 1.7522
Epoch [1/	20]		d_loss: 1.0073		g_loss: 1.3173
Epoch [1/	20]		d_loss: 1.1840		g_loss: 2.1046
Epoch [1/	20]		d_loss: 1.0700		g_loss: 0.7278
Epoch [1/	20]		d_loss: 0.9186		g_loss: 1.5632
Epoch [1/	20]		d_loss: 0.9053		g_loss: 1.1757
Epoch [2/	20]		d_loss: 1.0057		g_loss: 0.8216
Epoch [2/	20]		d_loss: 0.9260		g_loss: 1.2068
Epoch [2/	20]		d_loss: 1.0479		g_loss: 1.5992
Epoch [2/	20]		d_loss: 1.1101		g_loss: 0.9789
Epoch [2/	20]		d_loss: 1.1327		g_loss: 1.0723
Epoch [2/	20]		d_loss: 1.2674		g_loss: 0.4947
Epoch [2/	20]		d_loss: 1.3174		g_loss: 1.0939
Epoch [2/	20]		d_loss: 0.7901		g_loss: 1.2068
Epoch [2/	20]		d_loss: 1.0003		g_loss: 1.1316

Epoch [2/	20]	d_loss: 1.1252	g_loss: 2.1196
Epoch [2/	20]	d_loss: 1.2768	g_loss: 2.0847
Epoch [2/	20]	d_loss: 1.2234	g_loss: 0.9241
Epoch [2/	20]	d_loss: 1.2168	g_loss: 0.9574
Epoch [2/	20]	d_loss: 0.9800	g_loss: 1.1470
Epoch [2/	20]	d_loss: 0.9704	g_loss: 1.0632
Epoch [3/	20]	d_loss: 1.2513	g_loss: 1.8542
Epoch [3/	20]	d_loss: 0.8739	g_loss: 1.4763
Epoch [3/	20]	d_loss: 0.9837	g_loss: 2.2136
Epoch [3/	20]	d_loss: 1.2459	g_loss: 1.4557
Epoch [3/	20]	d_loss: 1.1252	g_loss: 1.6854
Epoch [3/	20]	d_loss: 1.0313	g_loss: 0.8557
Epoch [3/	20]	d_loss: 1.0825	g_loss: 1.0836
Epoch [3/	20]	d_loss: 1.1981	g_loss: 1.8777
Epoch [3/	20]	d_loss: 1.0294	g_loss: 1.4330
Epoch [3/	20]	d_loss: 1.1590	g_loss: 1.5991
Epoch [3/	20]	d_loss: 1.0552	g_loss: 1.2931
Epoch [3/	20]	d_loss: 0.9492	g_loss: 1.1740
Epoch [3/	20]	d_loss: 0.9858	g_loss: 1.0084
Epoch [3/	20]	d_loss: 1.0645	g_loss: 1.4179
Epoch [3/	20]	d_loss: 0.9441	g_loss: 1.0812
Epoch [4/	20]	d_loss: 1.2621	g_loss: 1.5526
Epoch [4/	20]	d_loss: 1.0309	g_loss: 1.6425
Epoch [4/	20]	d_loss: 1.1003	g_loss: 0.8838
Epoch [4/	20]	d_loss: 1.0158	g_loss: 1.3945
Epoch [4/	20]	d_loss: 1.1909	g_loss: 0.6678
Epoch [4/	20]	d_loss: 0.9112	g_loss: 0.9705
Epoch [4/	20]	d_loss: 1.1023	g_loss: 1.3225
Epoch [4/	20]	d_loss: 0.9295	g_loss: 1.3371
Epoch [4/	20]	d_loss: 1.3349	g_loss: 1.2650
Epoch [4/	20]	d_loss: 1.0429	g_loss: 0.9364
Epoch [4/	20]	d_loss: 1.5581	g_loss: 2.6967
Epoch [4/	20]	d_loss: 0.8518	g_loss: 1.4420
Epoch [4/	20]	d_loss: 1.1658	g_loss: 0.9626
Epoch [4/	20]	d_loss: 1.0161	g_loss: 1.9128
Epoch [4/	20]	d_loss: 1.2549	g_loss: 1.6670
Epoch [5/	20]	d_loss: 1.0224	g_loss: 0.9585
Epoch [5/	20]	d_loss: 0.9697	g_loss: 1.3222
Epoch [5/	20]	d_loss: 0.9127	g_loss: 1.6396
Epoch [5/	20]	d_loss: 0.9980	g_loss: 0.9969
Epoch [5/	20]	d_loss: 1.3951	g_loss: 0.9512
Epoch [5/	20]	d_loss: 1.1119	g_loss: 1.2140
Epoch [5/	20]	d_loss: 1.1610	g_loss: 0.7973
Epoch [5/	20]	d_loss: 0.9650	g_loss: 1.4245
Epoch [5/	20]	d_loss: 1.3614	g_loss: 1.5490
Epoch [5/	20]	d_loss: 0.9901	g_loss: 1.0379
Epoch [5/	20]	d_loss: 1.0795	g_loss: 1.4740
Epoch [5/	20]	d_loss: 0.9378	g_loss: 0.9393

Epoch [5/	20]	d_loss: 1.0541	g_loss: 0.9625
Epoch [5/	20]	d_loss: 1.1215	g_loss: 1.7283
Epoch [5/	20]	d_loss: 0.9887	g_loss: 1.4018
Epoch [6/	20]	d_loss: 0.8670	g_loss: 1.1408
Epoch [6/	20]	d_loss: 1.0011	g_loss: 1.3362
Epoch [6/	20]	d_loss: 1.1036	g_loss: 1.3586
Epoch [6/	20]	d_loss: 0.8444	g_loss: 0.6748
Epoch [6/	20]	d_loss: 1.1064	g_loss: 0.5956
Epoch [6/	20]	d_loss: 0.9141	g_loss: 0.6877
Epoch [6/	20]	d_loss: 1.1119	g_loss: 0.9553
Epoch [6/	20]	d_loss: 1.0663	g_loss: 1.1732
Epoch [6/	20]	d_loss: 1.3755	g_loss: 2.0979
Epoch [6/	20]	d_loss: 0.8307	g_loss: 1.6909
Epoch [6/	20]	d_loss: 0.9496	g_loss: 1.3523
Epoch [6/	20]	d_loss: 0.8106	g_loss: 1.3695
Epoch [6/	20]	d_loss: 1.0070	g_loss: 0.8168
Epoch [6/	20]	d_loss: 0.7614	g_loss: 1.4107
Epoch [6/	20]	d_loss: 0.9773	g_loss: 1.3105
Epoch [7/	20]	d_loss: 0.8796	g_loss: 1.7589
Epoch [7/	20]	d_loss: 0.8442	g_loss: 1.7696
Epoch [7/	20]	d_loss: 0.9290	g_loss: 1.5166
Epoch [7/	20]	d_loss: 0.7971	g_loss: 1.2538
Epoch [7/	20]	d_loss: 1.3050	g_loss: 0.6180
Epoch [7/	20]	d_loss: 0.8229	g_loss: 1.0684
Epoch [7/	20]	d_loss: 0.8912	g_loss: 0.9604
Epoch [7/	20]	d_loss: 0.4972	g_loss: 1.9471
Epoch [7/	20]	d_loss: 0.7046	g_loss: 1.2223
Epoch [7/	20]	d_loss: 1.0304	g_loss: 0.9111
Epoch [7/	20]	d_loss: 0.8530	g_loss: 1.4095
Epoch [7/	20]	d_loss: 0.6777	g_loss: 1.9372
Epoch [7/	20]	d_loss: 0.8228	g_loss: 1.6562
Epoch [7/	20]	d_loss: 0.8129	g_loss: 2.1938
Epoch [7/	20]	d_loss: 0.7455	g_loss: 1.7060
Epoch [8/	20]	d_loss: 0.8210	g_loss: 1.1450
Epoch [8/	20]	d_loss: 0.8405	g_loss: 1.7259
Epoch [8/	20]	d_loss: 0.8548	g_loss: 1.6497
Epoch [8/	20]	d_loss: 0.6861	g_loss: 1.2953
Epoch [8/	20]	d_loss: 1.0859	g_loss: 0.5681
Epoch [8/	20]	d_loss: 0.9105	g_loss: 0.6603
Epoch [8/	20]	d_loss: 0.7669	g_loss: 1.0739
Epoch [8/	20]	d_loss: 0.6335	g_loss: 1.2467
Epoch [8/	20]	d_loss: 1.0648	g_loss: 2.6402
Epoch [8/	20]	d_loss: 0.6987	g_loss: 1.5866
Epoch [8/	20]	d_loss: 0.9201	g_loss: 1.4020
Epoch [8/	20]	d_loss: 0.7432	g_loss: 1.0046
Epoch [8/	20]	d_loss: 0.5945	g_loss: 1.4572
Epoch [8/	20]	d_loss: 0.7987	g_loss: 2.1379
Epoch [8/	20]	d_loss: 0.9306	g_loss: 2.4405

Epoch [9/	20]	d_loss: 0.8540	g_loss: 1.4815
Epoch [9/	20]	d_loss: 0.8943	g_loss: 2.1157
Epoch [9/	20]	d_loss: 0.8691	g_loss: 1.3261
Epoch [9/	20]	d_loss: 0.7507	g_loss: 1.1041
Epoch [9/	20]	d_loss: 1.8466	g_loss: 0.4438
Epoch [9/	20]	d_loss: 0.7156	g_loss: 1.5514
Epoch [9/	20]	d_loss: 0.6838	g_loss: 1.9572
Epoch [9/	20]	d_loss: 0.4504	g_loss: 1.9651
Epoch [9/	20]	d_loss: 0.6725	g_loss: 0.9480
Epoch [9/	20]	d_loss: 0.8477	g_loss: 2.2960
Epoch [9/	20]	d_loss: 0.7440	g_loss: 1.1142
Epoch [9/	20]	d_loss: 0.7324	g_loss: 2.6842
Epoch [9/	20]	d_loss: 0.9042	g_loss: 1.9172
Epoch [9/	20]	d_loss: 0.8817	g_loss: 2.0246
Epoch [9/	20]	d_loss: 0.7863	g_loss: 1.8130
Epoch [10/	20]	d_loss: 0.8587	g_loss: 1.7674
Epoch [10/	20]	d_loss: 0.9957	g_loss: 1.8036
Epoch [10/	20]	d_loss: 0.6021	g_loss: 1.5647
Epoch [10/	20]	d_loss: 0.7990	g_loss: 2.5954
Epoch [10/	20]	d_loss: 0.9840	g_loss: 1.2172
Epoch [10/	20]	d_loss: 0.5301	g_loss: 1.4668
Epoch [10/	20]	d_loss: 0.6544	g_loss: 1.2098
Epoch [10/	20]	d_loss: 0.3954	g_loss: 2.7319
Epoch [10/	20]	d_loss: 0.6500	g_loss: 1.3433
Epoch [10/	20]	d_loss: 0.6584	g_loss: 2.6642
Epoch [10/	20]	d_loss: 0.8671	g_loss: 1.7520
Epoch [10/	20]	d_loss: 0.5751	g_loss: 2.2429
Epoch [10/	20]	d_loss: 0.5664	g_loss: 2.7798
Epoch [10/	20]	d_loss: 0.6543	g_loss: 1.5736
Epoch [10/	20]	d_loss: 0.7166	g_loss: 0.9169
Epoch [11/	20]	d_loss: 1.1040	g_loss: 2.8111
Epoch [11/	20]	d_loss: 0.5921	g_loss: 1.5074
Epoch [11/	20]	d_loss: 0.5983	g_loss: 1.9776
Epoch [11/	20]	d_loss: 0.6688	g_loss: 1.1112
Epoch [11/	20]	d_loss: 0.9944	g_loss: 1.3872
Epoch [11/	20]	d_loss: 0.5391	g_loss: 2.3324
Epoch [11/	20]	d_loss: 0.6676	g_loss: 2.0671
Epoch [11/	20]	d_loss: 0.6297	g_loss: 1.7209
Epoch [11/	20]	d_loss: 1.2555	g_loss: 3.3177
Epoch [11/	20]	d_loss: 0.8165	g_loss: 3.5583
Epoch [11/	20]	d_loss: 0.5738	g_loss: 1.3441
Epoch [11/	20]	d_loss: 0.5576	g_loss: 2.4064
Epoch [11/	20]	d_loss: 0.5270	g_loss: 1.4145
Epoch [11/	20]	d_loss: 0.5709	g_loss: 1.8934
Epoch [11/	20]	d_loss: 0.9731	g_loss: 2.9229
Epoch [12/	20]	d_loss: 0.9454	g_loss: 1.1146
Epoch [12/	20]	d_loss: 0.5293	g_loss: 2.3767
Epoch [12/	20]	d_loss: 0.4577	g_loss: 1.6173

Epoch [12/	20]	d_loss: 0.5427	g_loss: 1.4374
Epoch [12/	20]	d_loss: 0.7758	g_loss: 0.6743
Epoch [12/	20]	d_loss: 1.1891	g_loss: 2.6057
Epoch [12/	20]	d_loss: 0.4601	g_loss: 1.6892
Epoch [12/	20]	d_loss: 0.5554	g_loss: 1.8517
Epoch [12/	20]	d_loss: 0.5846	g_loss: 1.7792
Epoch [12/	20]	d_loss: 0.6275	g_loss: 2.0653
Epoch [12/	20]	d_loss: 0.7726	g_loss: 1.8965
Epoch [12/	20]	d_loss: 0.5542	g_loss: 1.9118
Epoch [12/	20]	d_loss: 0.5655	g_loss: 1.7891
Epoch [12/	20]	d_loss: 0.4890	g_loss: 1.3625
Epoch [12/	20]	d_loss: 0.6606	g_loss: 2.1471
Epoch [13/	20]	d_loss: 0.5435	g_loss: 1.7101
Epoch [13/	20]	d_loss: 0.6530	g_loss: 2.1914
Epoch [13/	20]	d_loss: 0.6449	g_loss: 1.5639
Epoch [13/	20]	d_loss: 0.6654	g_loss: 2.6487
Epoch [13/	20]	d_loss: 1.0104	g_loss: 1.8117
Epoch [13/	20]	d_loss: 0.4323	g_loss: 1.5480
Epoch [13/	20]	d_loss: 0.4679	g_loss: 2.0748
Epoch [13/	20]	d_loss: 0.3135	g_loss: 1.3723
Epoch [13/	20]	d_loss: 0.4980	g_loss: 1.5192
Epoch [13/	20]	d_loss: 1.1467	g_loss: 3.8900
Epoch [13/	20]	d_loss: 0.6293	g_loss: 1.4828
Epoch [13/	20]	d_loss: 0.9025	g_loss: 1.0082
Epoch [13/	20]	d_loss: 0.4315	g_loss: 1.7448
Epoch [13/	20]	d_loss: 0.5079	g_loss: 1.8569
Epoch [13/	20]	d_loss: 0.5872	g_loss: 2.3551
Epoch [14/	20]	d_loss: 0.7500	g_loss: 1.7908
Epoch [14/	20]	d_loss: 0.6992	g_loss: 2.4870
Epoch [14/	20]	d_loss: 0.4824	g_loss: 3.0691
Epoch [14/	20]	d_loss: 0.6089	g_loss: 2.8987
Epoch [14/	20]	d_loss: 0.8675	g_loss: 0.7803
Epoch [14/	20]	d_loss: 0.5113	g_loss: 0.8475
Epoch [14/	20]	d_loss: 0.4796	g_loss: 2.1828
Epoch [14/	20]	d_loss: 0.4331	g_loss: 2.8302
Epoch [14/	20]	d_loss: 0.4958	g_loss: 1.2750
Epoch [14/	20]	d_loss: 0.4831	g_loss: 1.6834
Epoch [14/	20]	d_loss: 0.6406	g_loss: 1.3663
Epoch [14/	20]	d_loss: 0.4392	g_loss: 2.4227
Epoch [14/	20]	d_loss: 0.4749	g_loss: 2.5893
Epoch [14/	20]	d_loss: 0.5727	g_loss: 2.0081
Epoch [14/	20]	d_loss: 0.6909	g_loss: 3.3945
Epoch [15/	20]	d_loss: 0.4974	g_loss: 2.5697
Epoch [15/	20]	d_loss: 0.3827	g_loss: 2.4618
Epoch [15/	20]	d_loss: 0.8037	g_loss: 1.9015
Epoch [15/	20]	d_loss: 0.3503	g_loss: 0.9107
Epoch [15/	20]	d_loss: 0.8069	g_loss: 1.1518
Epoch [15/	20]	d_loss: 0.4530	g_loss: 2.5920

Epoch [15/	20]	d_loss: 0.4420	g_loss: 2.5245
Epoch [15/	20]	d_loss: 0.3417	g_loss: 3.1737
Epoch [15/	20]	d_loss: 0.4439	g_loss: 1.4100
Epoch [15/	20]	d_loss: 0.3594	g_loss: 2.0192
Epoch [15/	20]	d_loss: 0.7240	g_loss: 1.6177
Epoch [15/	20]	d_loss: 0.3615	g_loss: 2.2966
Epoch [15/	20]	d_loss: 0.2967	g_loss: 2.4527
Epoch [15/	20]	d_loss: 0.3108	g_loss: 2.2035
Epoch [15/	20]	d_loss: 0.4708	g_loss: 2.0773
Epoch [16/	20]	d_loss: 0.4967	g_loss: 1.6103
Epoch [16/	20]	d_loss: 0.4905	g_loss: 2.4851
Epoch [16/	20]	d_loss: 0.3964	g_loss: 2.1665
Epoch [16/	20]	d_loss: 0.2904	g_loss: 3.0102
Epoch [16/	20]	d_loss: 1.1440	g_loss: 0.9409
Epoch [16/	20]	d_loss: 0.3308	g_loss: 2.0421
Epoch [16/	20]	d_loss: 0.3962	g_loss: 1.9286
Epoch [16/	20]	d_loss: 0.3437	g_loss: 2.8996
Epoch [16/	20]	d_loss: 0.8635	g_loss: 3.1621
Epoch [16/	20]	d_loss: 0.3686	g_loss: 2.2134
Epoch [16/	20]	d_loss: 0.4420	g_loss: 2.1367
Epoch [16/	20]	d_loss: 0.3622	g_loss: 2.9702
Epoch [16/	20]	d_loss: 0.4797	g_loss: 2.1969
Epoch [16/	20]	d_loss: 0.2184	g_loss: 2.0611
Epoch [16/	20]	d_loss: 0.4389	g_loss: 3.1045
Epoch [17/	20]	d_loss: 0.4218	g_loss: 1.6837
Epoch [17/	20]	d_loss: 0.4157	g_loss: 2.5578
Epoch [17/	20]	d_loss: 0.4740	g_loss: 2.1637
Epoch [17/	20]	d_loss: 0.2930	g_loss: 2.5534
Epoch [17/	20]	d_loss: 0.9577	g_loss: 1.4192
Epoch [17/	20]	d_loss: 0.4440	g_loss: 2.6027
Epoch [17/	20]	d_loss: 0.5257	g_loss: 2.0481
Epoch [17/	20]	d_loss: 0.2598	g_loss: 2.7104
Epoch [17/	20]	d_loss: 0.4821	g_loss: 1.2132
Epoch [17/	20]	d_loss: 0.2953	g_loss: 2.2692
Epoch [17/	20]	d_loss: 0.4920	g_loss: 0.8810
Epoch [17/	20]	d_loss: 0.6228	g_loss: 2.1119
Epoch [17/	20]	d_loss: 0.3434	g_loss: 2.0269
Epoch [17/	20]	d_loss: 0.2329	g_loss: 2.5180
Epoch [17/	20]	d_loss: 0.4720	g_loss: 2.2222
Epoch [18/	20]	d_loss: 0.4034	g_loss: 1.6099
Epoch [18/	20]	d_loss: 0.2509	g_loss: 3.0507
Epoch [18/	20]	d_loss: 0.2563	g_loss: 3.2339
Epoch [18/	20]	d_loss: 0.3215	g_loss: 3.6712
Epoch [18/	20]	d_loss: 0.6989	g_loss: 1.2042
Epoch [18/	20]	d_loss: 0.3903	g_loss: 3.0697
Epoch [18/	20]	d_loss: 0.4079	g_loss: 2.5796
Epoch [18/	20]	d_loss: 0.3742	g_loss: 3.0958
Epoch [18/	20]	d_loss: 0.4286	g_loss: 1.8928

```

Epoch [ 18/ 20] | d_loss: 0.2898 | g_loss: 2.9517
Epoch [ 18/ 20] | d_loss: 0.3272 | g_loss: 1.4707
Epoch [ 18/ 20] | d_loss: 0.2091 | g_loss: 2.9885
Epoch [ 18/ 20] | d_loss: 0.4379 | g_loss: 2.4095
Epoch [ 18/ 20] | d_loss: 0.3471 | g_loss: 2.1990
Epoch [ 18/ 20] | d_loss: 0.4438 | g_loss: 2.4786
Epoch [ 19/ 20] | d_loss: 0.6878 | g_loss: 2.8637
Epoch [ 19/ 20] | d_loss: 0.4942 | g_loss: 3.5926
Epoch [ 19/ 20] | d_loss: 0.2152 | g_loss: 2.0616
Epoch [ 19/ 20] | d_loss: 0.3006 | g_loss: 2.7155
Epoch [ 19/ 20] | d_loss: 0.8666 | g_loss: 0.9230
Epoch [ 19/ 20] | d_loss: 0.3419 | g_loss: 3.3136
Epoch [ 19/ 20] | d_loss: 0.5593 | g_loss: 2.6553
Epoch [ 19/ 20] | d_loss: 0.2126 | g_loss: 2.8030
Epoch [ 19/ 20] | d_loss: 0.3593 | g_loss: 1.2921
Epoch [ 19/ 20] | d_loss: 0.5482 | g_loss: 1.9387
Epoch [ 19/ 20] | d_loss: 0.3728 | g_loss: 2.1262
Epoch [ 19/ 20] | d_loss: 0.1574 | g_loss: 2.4795
Epoch [ 19/ 20] | d_loss: 0.2342 | g_loss: 2.7301
Epoch [ 19/ 20] | d_loss: 0.2543 | g_loss: 2.4874
Epoch [ 19/ 20] | d_loss: 0.2847 | g_loss: 1.9708
Epoch [ 20/ 20] | d_loss: 0.3749 | g_loss: 1.6547
Epoch [ 20/ 20] | d_loss: 0.5409 | g_loss: 3.0168
Epoch [ 20/ 20] | d_loss: 0.1401 | g_loss: 3.0096
Epoch [ 20/ 20] | d_loss: 0.3265 | g_loss: 3.4995
Epoch [ 20/ 20] | d_loss: 1.1819 | g_loss: 0.3874
Epoch [ 20/ 20] | d_loss: 0.4109 | g_loss: 3.2515
Epoch [ 20/ 20] | d_loss: 0.1746 | g_loss: 3.4196
Epoch [ 20/ 20] | d_loss: 0.2000 | g_loss: 2.6013
Epoch [ 20/ 20] | d_loss: 0.3369 | g_loss: 1.6194
Epoch [ 20/ 20] | d_loss: 0.7840 | g_loss: 2.9401
Epoch [ 20/ 20] | d_loss: 0.3343 | g_loss: 1.8324
Epoch [ 20/ 20] | d_loss: 0.1618 | g_loss: 2.7871
Epoch [ 20/ 20] | d_loss: 0.2252 | g_loss: 2.5882
Epoch [ 20/ 20] | d_loss: 0.4330 | g_loss: 4.1871
Epoch [ 20/ 20] | d_loss: 0.4243 | g_loss: 4.0376

```

2.8 Training loss


Plot the training losses for the generator and discriminator, recorded after each epoch.

```

In [19]: fig, ax = plt.subplots()
         losses = np.array(losses)
         plt.plot(losses.T[0], label='Discriminator', alpha=0.5)
         plt.plot(losses.T[1], label='Generator', alpha=0.5)
         plt.title("Training Losses")
         plt.legend()

```

Out[19]: <matplotlib.legend.Legend at 0x7f43783c55f8>




2.9 Generator samples from training

View samples of images from the generator, and answer a question about the strengths and weaknesses of your trained models.

```
In [20]: # helper function for viewing a list of passed in sample images
def view_samples(epoch, samples):
    fig, axes = plt.subplots(figsize=(16,4), nrows=2, ncols=8, sharey=True, sharex=True)
    for ax, img in zip(axes.flatten(), samples[epoch]):
        img = img.detach().cpu().numpy()
        img = np.transpose(img, (1, 2, 0))
        img = ((img + 1)*255 / (2)).astype(np.uint8)
        ax.xaxis.set_visible(False)
        ax.yaxis.set_visible(False)
        im = ax.imshow(img.reshape((32,32,3)))

In [21]: # Load samples from generator, taken while training
with open('train_samples.pkl', 'rb') as f:
    samples = pkl.load(f)

In [22]: _ = view_samples(-1, samples)
```



2.9.1 Question: What do you notice about your generated samples and how might you improve this model?

When you answer this question, consider the following factors: * The dataset is biased; it is made of "celebrity" faces that are mostly white * Model size; larger models have the opportunity to learn more features in a data feature space * Optimization strategy; optimizers and number of epochs affect your final result

Answer: I noticed that most of the generated images are biased (white faces). The model I made is very simple, I think having more variant data, deeper model, and more number of epochs will improve the result of the model. Furthermore, I noticed that when I train the model the Generator network loss is increasing at some point, while the Discriminator gets better and better, I'm not sure if this is a correct result.

2.9.2 Submitting This Project

When submitting this project, make sure to run all the cells before saving the notebook. Save the notebook file as "dlnd_face_generation.ipynb" and save it as a HTML file under "File" -> "Download as". Include the "problem_unittests.py" files in your submission.